



Beacon AI

Cross-border trend intelligence for
Indian venture capital — what is forming
in the US and China, scored for *transfer
to India*.


Working MVP

Venture Capitalists

Minimal risk, EU AI Act


The problem

A fund hears about a category once its Series A is priced — sometimes only at IPO.



\$0.93B VC software market, 2025 ~11.1% CAGR → \$1.58B by 2030 200+ Indian small-VCs: 4 to 10 staff, 1 research analyst	~4,500 Active VC firms globally the entire buyer universe	13→28% Use AI inside the decision doubled in one year (Affinity) Why it persists: the signal sits in thousands of foreign-language filings	600+ Tools already competing but ~20 used by 80% of firms
--	--	--	--

The lag, wave by wave



China's peak leads India's by one to two years; the US lead runs three to seven.

Since Round 1



What held

- The industry: VC software
- The buyer: **Indian VCs or VCs investing in India**
- The thesis remain unchanged: , priced out of Western platforms, **buying an early read on trends** forming in the US and China.

What changed

- Narrower and deeper trend framework
- Dataset list is trimmed to sources we can defend
- n8n orchestration layer is gone
- In their place is one working analyst application — the thing the fund actually opens

How Beacon AI works — and it works today



- Running today: retrieval, clustering and sourced generation, end to end
- Every claim carries its source; a partner decides what enters diligence
- Public, company-level data only — no personal data in the pipeline
- Production adds licensed sources, an evaluation harness and audit logs
- Runs in a browser — nothing to install

Pricing

ENTRY

\$300–500

per month

Solo GP or very small fund — the pilot price.

- ✓ Weekly transfer brief
- ✓ One seat
- ✓ One sector tracked
- ✓ PDF export

RECOMMENDED

CORE

\$6–12K

per year

A small fund team — the tier the ROI model uses.

- ✓ Everything in Entry, plus:
- ✓ Three seats
- ✓ Full sector coverage
- ✓ Custom watchlists
- ✓ Analysis on demand

GROWTH

\$18–30K

per year

Multi-seat funds, Beacon in the weekly process.

- ✓ Everything in Core, plus:
- ✓ Ten seats
- ✓ Category alerts
- ✓ Workflow integration
- ✓ Quarterly review call

ENTERPRISE


Custom

annual, on scope

Larger funds or broader private-markets use.

- ✓ Everything in Growth, plus:
- ✓ API access
- ✓ DPA and SSO
- ✓ Client-specified sources
- ✓ Dedicated analyst hours

Costs, value and ROI



Cost barely moves as we grow — each additional fund costs under \$150 a month to serve.
Estimates on industry benchmarks, validated in the pilot.

The risk register

RISK	CATEGORY	LIKELIHOOD × IMPACT	OWNER
Unsupported claim in a brief	Technical	Likely × High	Engineering
Transfer thesis fails for a category	Technical	Possible × High	Research
Micro-VCs will not pay at the core tier	Commercial	Likely × Medium	GTM
Model or search provider changes terms	Operational	Possible × Medium	Engineering
Source licensing dispute over reports	Regulatory	Possible × High	Legal counsel
Scope creep to person-level scoring	Regulatory	Unlikely × High	Compliance
Herding — subscribers chase one category	Ethical	Possible × Medium	Research
Key-person dependency on one builder	Operational	Likely × High	Founder

Each risk carries a mitigation and a review date; the first three are my focus.

Compliance

EU AI Act — minimal risk

- Outside Annex III: no decision about a natural person, no biometrics, no emotion inference. The subject is companies.
- What applies: Article 4 AI literacy, Article 50 disclosure, deployer duties on a general-purpose model.
- No conformity assessment or CE (Conformité Européenne) marking at this class.


GDPR and data flows

- Signal pipeline is company-level by design; no person-level profile is built. Personal data is limited to subscriber accounts.
- Inference is a US sub-processor under a DPA with standard contractual clauses, zero-retention, no personal data in prompts.
- A full DPIA is not mandated; the screening DPIA capped retention at thirty days.

Reclassification in future: The day Beacon AI scores individuals, or feeds a credit or hiring decision, it is re-assessed before shipping


BEACON · EXECUTION PLAN

Sprint view — three phases, ten sprints



Every story also meets a shared 7-point Definition of Done, has detailed scope and acceptance criteria — see the reference slide.

PoC to MVP to launch



60%

of briefs rated
actionable by partners

3 of 5

design partners
convert to paid

≤5%

factual-error rate on
sampled claims

\$6–12K

core annual seat price
at launch



A 60-day pilot

Three to five Indian micro-VCs, weekly briefs, four success criteria agreed up front.

Beacon AI turns a 12-to-36-month category lag into proprietary deal flow, at a price a \$10–40M fund can carry. Minimal risk under the EU AI Act, company-level data by design. The decision in front of you is the pilot.


*“You are saving my time
but if you can replace
my analyst, you have
my money.”*

Managing Director and Partner, VC
(India focus)



APPENDIX

Reference material



Costs and ROI in detail

Cost to build — \$28K

Pilot-ready product: ~\$3K cash for LLM compute and other API licences, plus development and consulting

Value to the fund

About \$20K a year of analyst time returned, against a \$6–12K subscription — and one earlier entry into one winning category that 'could do an IPO'.

Estimates on industry benchmarks — validated in the pilot

Cost to run — \$18K a year

At ten funds:

- \$11K inference and search, \$3K cloud and storage, \$4K enrichment and support
- Marginal cost per additional fund is under \$150 a month.

ROI — +57% at 12 months

- Eight funds at ~\$9K average: ~\$72K against ~\$46K of cost.
- At 36 months, ~25 funds and ~+240%.
- Break-even at six funds, month nine

KPIs & acceptance criteria — PoC → MVP

The stitch, sprint by sprint: each sprint’s stories, the KPIs they’re measured by, and the acceptance criteria that close them.

SPRINTS 0–1 · POC

Cross-Border Funding Data Foundation · Transfer-Probability Signal Engine · Conversational Front Door (n8n)

KPIS · PER STORY

- Time from chat message sent → brief returned in chat (baseline to be measured and recorded, since no target currently exists in the repo).
- Brief generation success rate (runs that return a valid structured brief vs. fail/timeout).

ACCEPTANCE CRITERIA · PER SPRINT

- Project setup completed, folder structure defined
- Files contain ranked sectors with scores, restricted to sectors with nonzero China/US deal presence
- Files contain a generated brief per top-ranked sector plus a combined brief
- End-to-end chat demo: a partner names a valid sector and receives a brief in-chat; an invalid/no-match input returns the sector picklist instead of an error
- At least one generated run is visible as a trace in LangSmith

SPRINT 2 · MVP

Trend Hierarchy Generation · Live Research Grounding

KPIS · PER STORY

- % of runs completing with live research coverage vs. falling back to sample data.
- Run completion rate without manual intervention.

ACCEPTANCE CRITERIA · PER SPRINT

- App takes user inputs and passes to the engine
- A run with all four inputs set returns Macro/Mega/Sub-trend cards, each containing every required field
- All unit tests pass, confirming Serper → SerpAPI → Tavily fallback order
- Research context is written on every run with provider order, live-hit count, and extracted keywords populated
- A forced truncated-response condition falls back to sample data instead of crashing or rendering partial output

SPRINT 3 · MVP

Momentum & Trend Radar Visualization · News Signals

KPIS · PER STORY

- News relevance rate (articles matching selected industry + region, sampled manually per release since no automated relevance scorer exists yet).

ACCEPTANCE CRITERIA · PER SPRINT

- Trend Radar chart renders with an explicit title (no “undefined”) and the strongest category visually distinguished
- News Signals tab renders exactly 8 articles per run
- No article is off-industry or off-region
- News fallback to Currents triggers correctly when NewsAPI returns fewer than 5 results or does not yield any result

SPRINT 4 · MVP

Sub-Trend Drill-Down · Industry Auto-Detection

KPIS · PER STORY

- Industry-detection eval accuracy score against the ground-truth set (target to be set once a baseline run exists).

ACCEPTANCE CRITERIA · PER SPRINT

- Drill-down for a sub-trend returns companies and social signals scoped to the selected time range; trending products are either industry-matched or an explicit empty state — never an unfiltered, irrelevant fallback list.
- Website→industry detector eval run against the ground-truth set, with accuracy/honesty/helpfulness scores recorded in `langsmith/langsmith_documentation.md`.

KPIs & acceptance criteria — MVP → Launch

Continuing the stitch through the final sprints, plus the Definition of Done every story must meet.

SPRINT 5 • MVP

Summary Report & Export

KPIS • PER STORY

- Share / download success rate.

ACCEPTANCE CRITERIA • PER SPRINT

- Summary Report tab renders the combined markdown + trend table for a completed run.
- PDF download reads from the saved file on disk under `raw/analysis/`, not from in-memory state only; the share link is shown only on confirmed successful upload — a failed upload disables sharing for that run without blocking the local PDF download.

SPRINT 6 • MVP

Reliability, Cost & Performance Engineering • Quality & Regression Safety

KPIS • PER STORY

- 100% of OpenAI call sites covered by failover.
- 100% of runs producing a cost log entry.

ACCEPTANCE CRITERIA • PER SPRINT

- All 6 OpenAI call sites use rotational keys.
- Every analysis run writes logs with token counts, latency, and estimated cost populated.
- First-click (“cold”) analysis latency is measured and recorded as a baseline metric, even if not yet at target.

SPRINTS 7–9 • LAUNCH

Onboarding flow for humans • Identify India-focussed companies analogous to the ones trending in US and China • Build a landscape of the companies operating in that particular sub-trend • Agent access

KPIS • PER STORY

- Multi-tenant isolation — 0 cross-account data leaks across test accounts (a hard requirement, not a target range, since any nonzero value is a data breach).
- % of trending non-India companies for which the app surfaces at least one India lookalike.
- Onboarding funnel — time from signup to first completed analysis run (baseline to be measured once tenant isolation exists).
- Agent/API-triggered run success rate (target to be set once a baseline run exists, per the pattern used for other not-yet-baselined KPIs).

ACCEPTANCE CRITERIA • PER SPRINT

- A second user can sign up, log in, and run an analysis without seeing another account’s runs, saved profile, or preferences — verified with two concurrent test accounts, not just single-user testing. Profile/preference values persist across a full logout/login cycle, not just within one browser session.
- For a completed analysis run, the app returns at least one India-based “lookalike” company for each trending China/US comparison company surfaced in that run, each tagged with funding stage, company size, and growth rate — never an unlabeled name with no comparison basis.
- The ecosystem landscape view for a selected sub-trend renders using real drill-down data (companies/signals already generated for that sub-trend), not placeholder or unrelated content; an empty state shows when no drill-down has been generated yet.
- An external caller can trigger a trend-generation run and retrieve the resulting hierarchy through an API call alone, with no browser/Streamlit session involved.

DEFINITION OF DONE

Applies to every story — no exceptions.

- 1 Code merged behind the feature’s owning module
- 2 A live path and a deterministic sample/mock fallback both exist and are exercised — never just the happy path
- 3 Automated tests exist for the story’s behavior and pass (pytest)
- 4 Manually verified end-to-end in the running app (not just unit-tested in isolation) before being called done
- 5 Failure modes handled visibly: a plain-language warning to the user, never a raw stack trace or a silent wrong answer
- 6 No secrets (API keys, tokens) committed; all credentials read from a local, gitignored `.env`
- 7 Behavior documented in the owning README — the provider order, validation rule, or fallback behavior the story introduces